



# CSCI-UA.0480-003

# Parallel Computing

## Lecture 2: Parallel Hardware: Basics


Mohamed Zahran (aka Z)

mzahran@cs.nyu.edu

<http://www.mzahran.com>



# Computer History




ENIAC

Eckert and Mauchly



- 1<sup>st</sup> working electronic computer (1946)
- 18,000 Vacuum tubes
- 1,800 instructions/sec
- 3,000 ft<sup>3</sup>

# Computer History



EDSAC 1 (1949)

<http://www.cl.cam.ac.uk/UoCCL/misc/EDSAC99/>

- Maurice Wilkes




1<sup>st</sup> stored program  
computer

650 instructions/sec

1,400 ft<sup>3</sup>


# Intel 4004 Die Photo



- Introduced in 1970
  - First microprocessor
- 2,250 transistors
- 12 mm<sup>2</sup>
- 108 KHz




# Intel 8086 Die Scan



- 29,000 transistors
- 33 mm<sup>2</sup>
- 5 MHz
- Introduced in 1979
  - Basic architecture of the IA32 PC

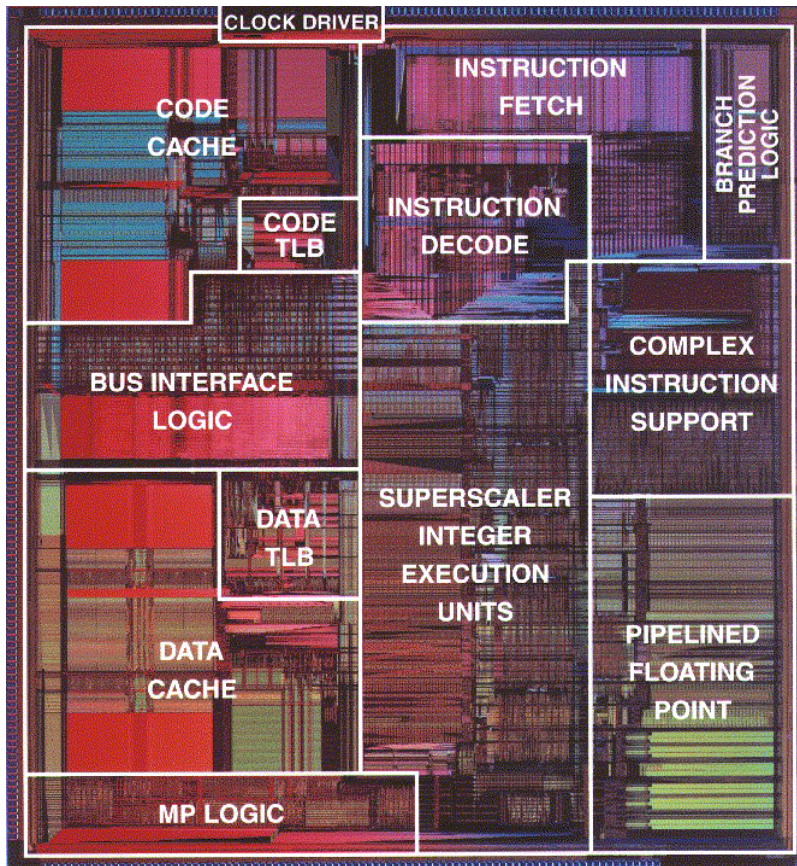


# Intel 80486 Die Scan




- 1,200,000 transistors
- 81 mm<sup>2</sup>
- 25 MHz
- Introduced in 1989
  - 1<sup>st</sup> pipelined implementation of IA32
  - 1<sup>st</sup> processor with on-chip cache

# Pentium Die Photo



- 3,100,000 transistors
- 296 mm<sup>2</sup>
- 60 MHz
- Introduced in 1993
  - 1<sup>st</sup> superscalar implementation of IA32

# Pentium 4



55,000,000  
transistors


146 mm<sup>2</sup>

3 GHz


Introduced in 2000

<http://www.chip-architect.com>






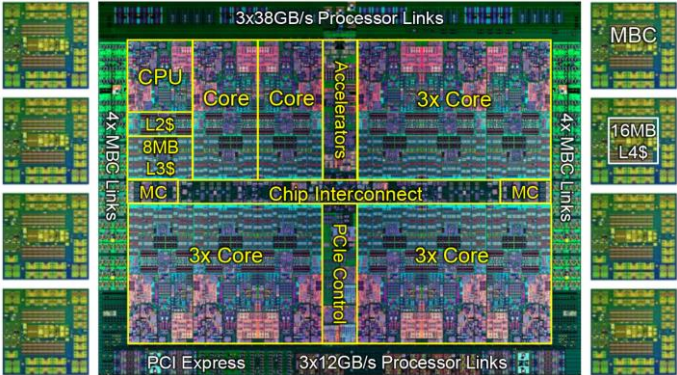
Pentium 4




Core 2 Duo (Merom)



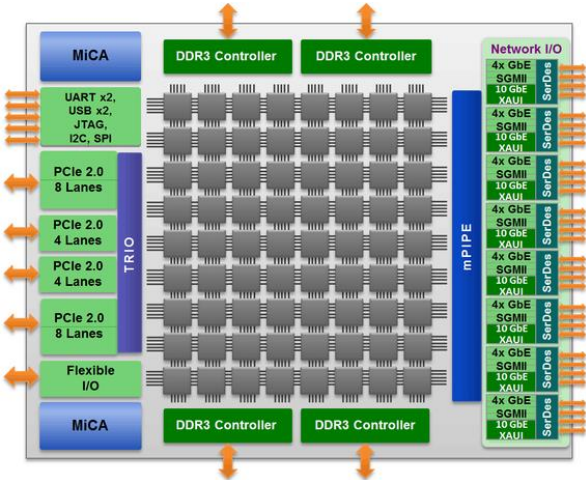
Intel Core i7 (Nehalem)



IBM Power 8




Intel Xeon Phi (50 cores)



Tiler (72 cores)

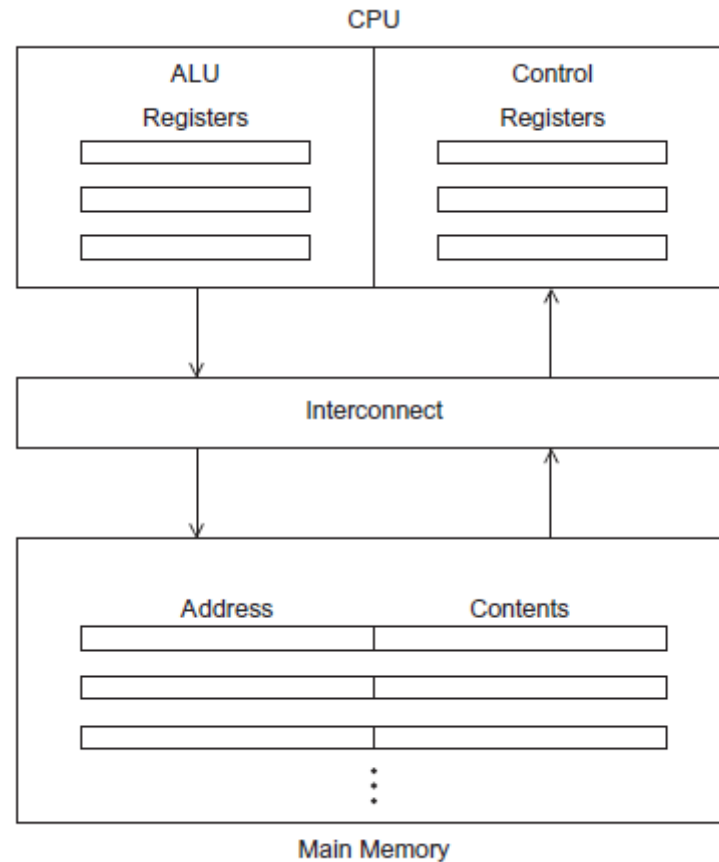
How did the hardware evolve like  
that?

# First Generation (1970s)




Single Cycle Implementation

# The Von Neumann Architecture





# Second Generation (1980s)



- **Pipelining:**

- the hardware divided into stages
  - **temporal parallelism**
  - Number of stages increases with each generation
- 
- Maximum **CPI** (Cycles Per Instruction) = 1
    - Due to dependencies  
(i.e. an instruction must wait for the result of another instruction)

# Some Enhancements

```
graph TD; A[Some Enhancements] --> B[Cache Memory]; A --> C[Virtual Memory]; B --> D[Multi-level caches]; C --> E[TLB]
```

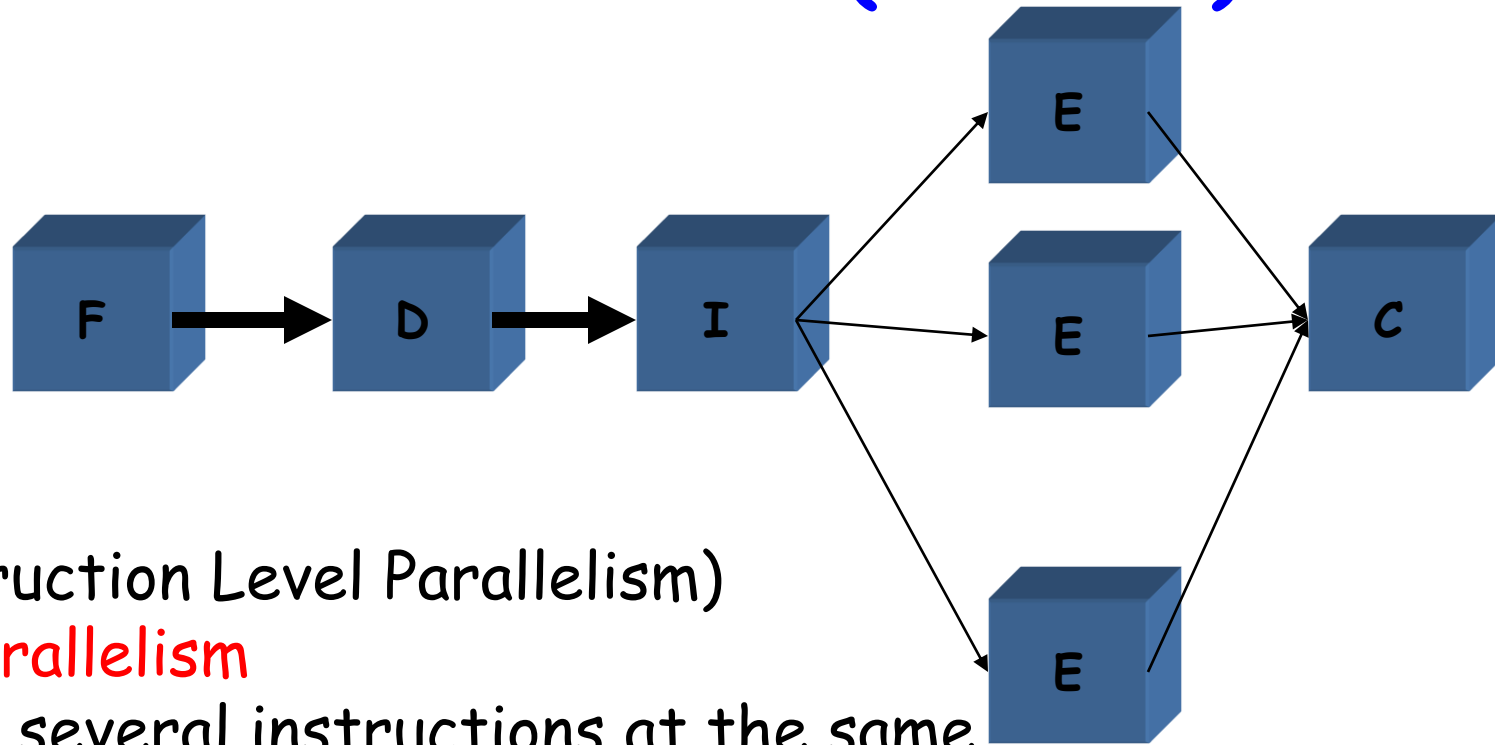
Cache Memory

Virtual Memory

Multi-level caches


TLB

# Third Generation (1990s)



- **ILP** (Instruction Level Parallelism)
- **Spatial parallelism**
- Executing several instructions at the same time is called **superscalar** capability.
- performance = instructions per cycle (**IPC**)
- **Speculative Execution** (prediction of branch direction) is introduced to make the best use of superscalar capability → This can make some instructions execute **out-of-order**!!

# Fourth Generation (2000s)



Simultaneous Multithreading (SMT)  
(aka Hyperthreading Technology)



Some definitions before we proceed

# An operating system "process"

- An instance of a computer program that is being executed.
- Components of a process:
  - The executable machine language program
  - A block of memory
  - Descriptors of resources the OS has allocated to the process
  - Security information
  - Information about the state of the process

# Multitasking

- Gives the illusion that a single processor system is running multiple programs simultaneously.
- Each process takes turns running → **time slice**
- After its time is up, it waits until it has a turn again.

# Threading

- Threads are **contained within processes**.
- They allow programmers to divide their programs into (more or less) independent **tasks**.
- The hope is that when one thread blocks because it is waiting on a resource, another thread will have work to do and can run.



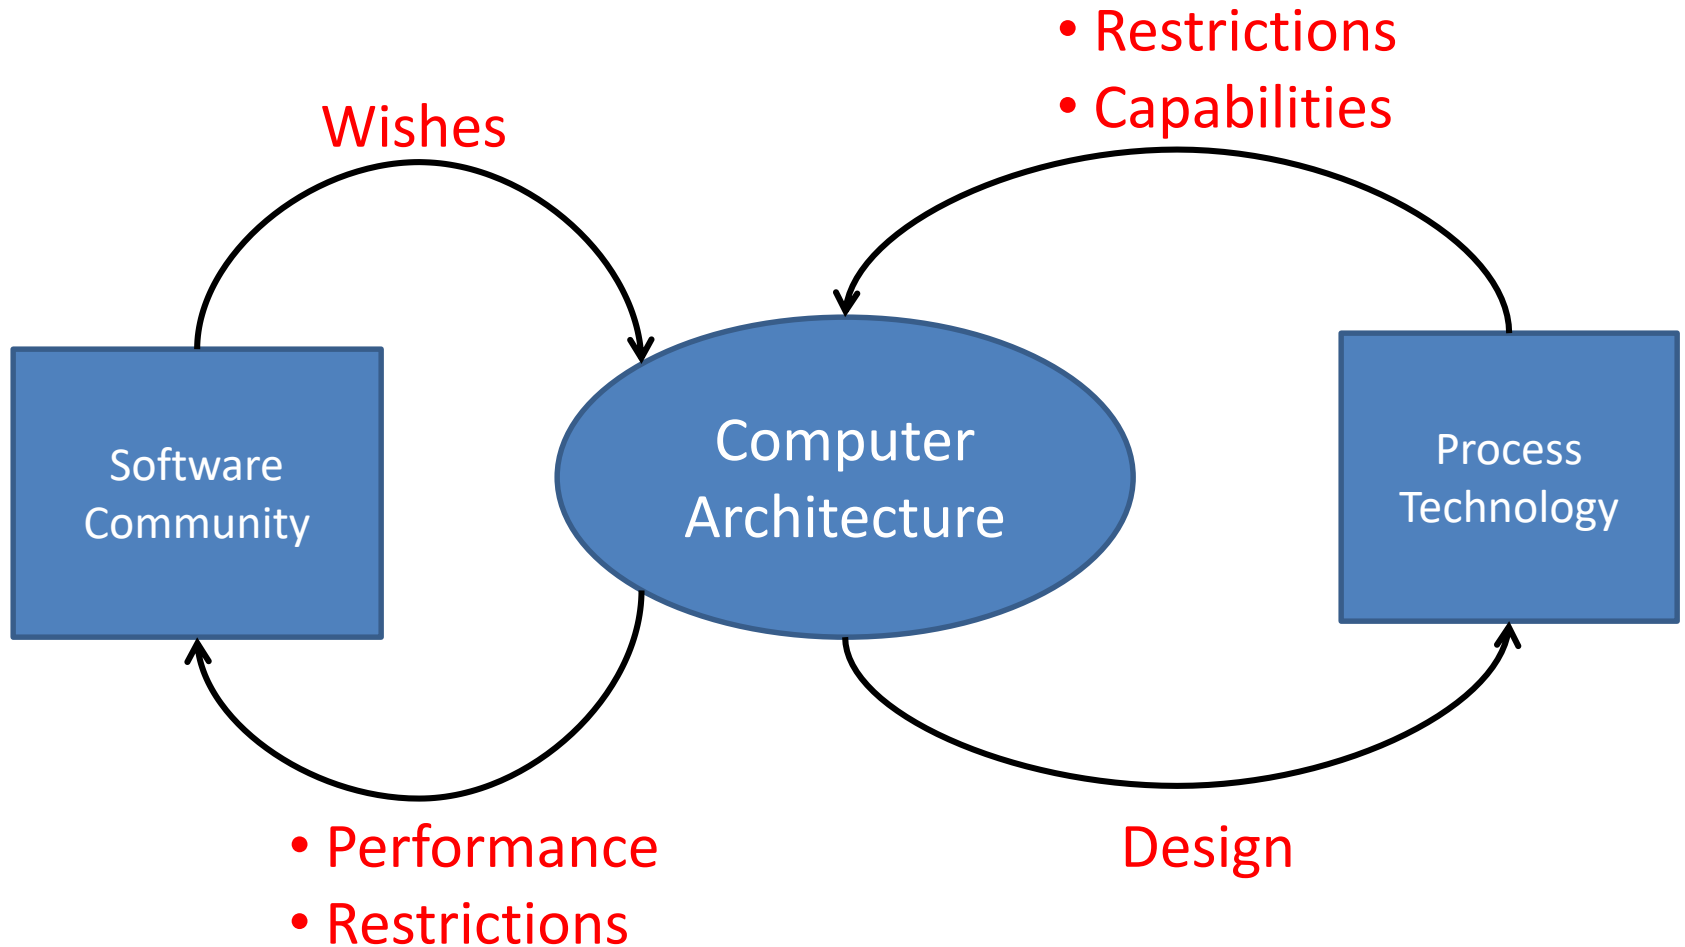
As you can see ...


We can have several processes, executed in a multitasking fashion, and each process can consist of several threads.

# The Status-Quo

- We moved from single core to multicore to manycore:
  - for technological reasons, as we saw last lecture.
- Free lunch is over for software folks
  - The software will not become faster with every new generation of processors
- Not enough experience in parallel programming
  - Parallel programs of old days were restricted to some elite applications -> very few programmers
  - Now we need parallel programs for many different applications

# How Did These Advances Happen?





The Multicore Software Triad



# Conclusions

- The hardware evolution, driven by Moore's law, was geared toward two things:
  - exploiting parallelism
  - Dealing with memory (latency, capacity)